

A Brief Haskell and GHC Refresher

Henrik Nilsson and Neil Sculthorpe
School of Computer Science
University of Nottingham

24th September 2013

1 Introduction

The purpose of this document is to give you quick recap of Haskell and GHC (the Haskell system we are using). This is recommended reading if you are taking the G52MAL or G53CMP modules. However it is there only for your benefit: it is not assessed and there is nothing to hand in. Should you feel you need a more in-depth Haskell refresher, or if you happen to be new to Haskell, Graham Hutton's book *Programming in Haskell* is highly recommended (in particular, the first 70 pages). Even if you are a fairly confident Haskell programmer, we advise you to quickly glance through this document: it may contain aspects of Haskell and GHC that you are not familiar with, but which are used in supplied code (particularly for G53CMP) as they are useful when developing large programs.

We recommend that you use the Haskell system GHC on the School's Linux servers. However, you could use other Haskell implementations, such as Hugs or NHC, or different platforms, such as Unix, Mac OS X, or Windows, if you prefer. In particular, the Haskell Platform (which includes GHC) has recently been installed on the Windows machines in the Lab. Note that if you use other platforms or Haskell implementations, then you cannot expect the course Teaching Assistants (TAs) to provide much technical assistance if you run into trouble with your installation. The site www.haskell.org is your starting point for everything you possibly want to know about Haskell, and for downloading Haskell implementations, related tools, and documentation.

2 Setting Up

2.1 Linux

The following assumes that you work on one of the School's Windows machines; e.g. in the main lab A32. Log on to your Linux server using e.g. the SSH Secure Shell Client (easiest) or PuTTY (there should be shortcuts to both on your desktop). At time of writing, the servers are *avon* for 1st year students, *bann* for 2nd year students and *clyde* for 3rd and 4th year students.

Start the interactive GHC environment by issuing the command `ghci` at the command line prompt:

```
bann$ ghci
```

Some information about GHCi gets printed, and you'll then get a new prompt:

```
Prelude>
```

From here, you can enter and evaluate Haskell expressions, load Haskell code from files, etc. See section 3 for more details.

You can also edit code on the servers using text editors like Emacs (command `emacs`) or Vi (command `vi`). Using a terminal multiplexer like Screen (command `screen`; do `man screen` for info) you can start a number of interactive sessions (e.g. GHCi, Emacs, shell) and quickly and easily switch between them, all within one window. Alternatively, you can start a number of SSH sessions in separate windows by invoking the SSH client multiple times.

2.2 Windows

The Haskell Platform, which includes GHCi, has been installed on the Windows machines in the lab. Just select GHCi from the start menu (you will find it under All Programs, Haskell Platform).

Note that you can navigate around the directory structure using the `:cd` command. For example, to get to the H drive:

```
:cd H:
```

Also, you can set GHCi (if it isn't already) as the default program associated with `.hs` files, so you can load them into GHCi just by clicking on them in a file browser window.

Alternatively, you can use WinGHCi. It allows you to do simple things like loading, editing, and running code through GUI shortcuts. However, the

associated editor is Notepad, and as Notepad does not understand Unix line-ending conventions, you may need to work around that one way or another in certain cases: see below.

You can edit Haskell files on the Windows machines using editors like XEmacs or Notepad++ (there should be shortcuts to both on your desktop). Both of these adapt automatically to different line-ending conventions, but Notepad++ may need some configuration regarding the width of tab stops: see below.

2.2.1 Unix and Windows Line-Ending Conventions

As you may be aware, Unix (and hence also Linux, Solaris, Mac OS, etc.) and Windows use different line-ending conventions for text files. Consequently, you could encounter problems if you switch between systems. In particular, the source code for the G53CMP coursework was created under Linux, and thus uses its line-ending convention. To get around this problem, you can use the `unix2dos` and `dos2unix` programs to convert text files from Unix to Windows and vice-versa (respectively). You can run these programs (under Linux) by supplying them with the names of one or more files to convert (old files will be overwritten); for example:

```
unix2dos MyFile1.hs MyFile2.hs MyFile3.hs
```

Alternatively, you can specify input-output file pairs; for example:

```
unix2dos -n MyFile-Unix.hs MyFile-Windows.hs
```

You can read about these issues in more detail in Wikipedia:

```
http://en.wikipedia.org/wiki/Newline
```

2.2.2 Haskell Layout and the Width of Tab Stops

Another issue concerns assumptions about the width of tab stops, although this is more of a tool issue (in particular, text editors, like Emacs or Notepad++) than an operating system issue.

Parsing of Haskell programs take layout (indentation) into account (unless the structure is made explicit using curly braces and semicolons). If tab characters are used in a Haskell file, it thus become a critical question just how wide (in spaces) a tab *stop* is supposed to be, as the presence of a tab character means that the horizontal position of the next character should be aligned with the next tab stop. The Haskell language standard has a precise definition (to ensure that Haskell programs always are interpreted the same

way): a tab stop is 8 spaces wide. This is also the default in many (most?) text editors, like Emacs.

However, for example Notepad++, which is a popular text editor among Windows users, has (at least usually) a different idea about the default: it opts for tab stops being 4 spaces wide. To avoid unnecessary grief caused by this (i.e., seemingly inexplicable parse errors), it is recommended that you, when editing Haskell source using Notepad++, go to Settings, Preferences, Tab Settings and change the width of tab stops to 8, and that you also tick the box “treat tabs like spaces”.

If using Notepad, at least from within WinGHCi, the width of a tab stop seems to default to 8, which is appropriate for Haskell, but as noted above, it seems Notepad cannot handle Unix line-ending conventions, so you might need to convert files from Unix to Windows conventions manually.

3 Getting Started

3.1 Using GHCi

The text before `>` (in this case `Prelude`) are the names of the modules whose definitions are in scope. Thus the above prompt means that everything in the Haskell standard Prelude is available.

Now evaluate some simple expressions. For example:

```
Prelude> 1 + 2
Prelude> "Hello World!"
Prelude> putStrLn "Hello World!"
Prelude> "Hello\nWorld!\n"
Prelude> putStrLn "Hello\nWorld!\n"
```

Make sure you become familiar with the command-line editing facilities to save on typing. For example, try out the arrow keys and various Emacs bindings.

Beside expressions, which get evaluated when entered, GHCi accepts a number of commands. They are all prefixed by `:` to distinguish them from Haskell expressions. The command `:help` prints a list of available commands. Try it now. Note that commands may be abbreviated to save on typing. For example `:h` is enough to get help. Try it. What is the command to load a Haskell module (and recursively all modules it depends on) from a file? What is the command for leaving GHCi? Try it now, and then restart GHCi.

3.2 Haskell Online

Visit `www.haskell.org` and locate:

1. The page on Haskell tutorials.
2. The book *Learn You a Haskell for Great Good!*
3. The GHC documentation, in particular the section on GHCi.

4 Trees

Task 1

Using a text editor of your choice, define a module called **Tree**. The GHC convention is that there should be one module per file, and that the name of the file should be the same as the name of the module defined in it with an additional `.hs` extension. (This is how GHC can find the definitions of a module given just its name.) In your case the file should be called **Tree.hs**.

The module **Tree** should contain a data declaration for a tree with three data constructors representing:

- an empty tree
- a singleton tree (a leaf)
- a non-empty tree consisting of two subtrees

Both the singleton tree node and the interior tree nodes should carry a single `Int` value. Call the constructors **Empty**, **Leaf**, and **Node**, for example.

Task 2

Load the module into GHCi. If there are errors, fix them and try again. Note how the prompt changes to indicate that the definitions in the module **Tree** now are in scope. Which are they? What are their types? (Hint: try the command `:type.`)

As the module **Prelude** is implicitly imported into every module unless explicitly hidden, all Prelude definitions are still available. (The `*` in the prompt `*Tree>` means that the Prelude is in scope.)

Task 3

Construct some `Tree` values. Can you print them? Can you compare them using the operators `==` or `<`? If not, fix the problem (hint: make use of a `deriving` declaration) and reload the module.

Task 4

The command `:show modules` shows all loaded modules. Try it. Switch back so that only Prelude definitions are in scope again. Command `:module Prelude` (or just `:m Prelude`). Are your `Tree` type and data constructors now available?

You can still get at your definitions by using their *fully qualified names*. That is, by prefixing the name of a defined entity with the name of the module in which it is defined. For example, if one of your `Tree` data constructors is called `Node`, it is available as `Tree.Node`. Try this. This works because GHCi as an extra convenience implicitly imports the definitions of all loaded modules under their fully qualified names into all scopes. (When a module is compiled by any Haskell 98 compiler, such as GHC, all used definitions from other modules must be explicitly imported in one way or another, except for those from the Prelude.) Now switch back to the `Tree` scope.

Task 5

Generalise the your `Tree` definition so that the tree can carry data of an arbitrary type. That is, make it *polymorphic*. Load the new definition into GHCi and test it. What are the types of the `Tree` data constructors now? Make sure you understand them!

Task 6

Using your polymorphic data constructors, create trees of integers, characters, and strings. Check that you can print and compare them and that they have the expected type.

Task 7

Create a new module called `Main` (in the file `Main.hs`). Import the module `Tree` into this module.

Task 8

Define a function `size` that returns the size of a tree. The size of the tree is the number of values carried by the tree.

Task 9

Load the new module into GHCi and test it on trees of a few different sizes. (Hint: you may want to define a number of test trees under some convenient names, either in the module `Main` or in a separate third module of containing test data that you import into `Main`.)

Check which modules are loaded now. Switch between the different module scopes and figure out which definitions are available where (without giving their fully qualified names).

Task 10

Define a function `insert` that inserts a value at the right place in an ordered tree. A tree is ordered if all values in the left sub tree is strictly smaller than the value in the top node, and all values in the right subtree is strictly greater than the value of the top node. A particular value occurs at most once in a tree.

Task 11

Load the new version of `Main`. What is the type of `insert`. Why? Add an explicit type signature to the definition of `insert` as documentation!

Task 12

Change your `Tree` definition so that it uses named (also called *labelled*) fields. This is Haskell's version of records. Let the name for all value fields be `value`. Let the names for the left and right subtree be `left` and `right` respectively.

Task 13

Load the new definition. Verify that you can construct trees using the named field notation, and that the order among the fields does not matter when you do so. For example, assuming the constructors are called `Empty`, `Leaf`, and `Node`:

```
*Main> Node {left=Leaf {value=1}, value=2, right=Empty}
*Main> Node {right=Empty, left=Leaf {value=1}, value=2}
```

Verify that you still can construct trees using the normal way of applying the constructor functions (i.e. with positional arguments). For example

```
*Main> Node (Leaf 1) 2 Empty
```

Note that the result tree still gets printed using the named field notation.

What happens if you don't provide values for all fields? Or indeed for no field? For example, what are the results of the following? Why?

```
*Main> :type Node {left=Empty}
*Main> :type Node {}
*Main> Node {left=Empty}
*Main> Node {}
```

Verify that you automatically got selector functions `value`, `left`, and `right`, that these have the expected types, and that you can use them to pick trees apart. What happens if you apply these selector functions to trees with the wrong top-level constructor, such as:

```
*Main> value Empty
*Main> left (Leaf 2)
```

Explain.

Task 14

Redefine your `size` and `insert` functions so that they make use of the field names when pattern matching. Verify that you can match the fields in any order. Also note that you easily can omit field names that are of no interest (e.g. `value` when you're defining the function `size`). It may also make sense to leave out field names that are only of interest in certain conditional branches of the code to make the pattern matching clearer and draw attention to what is most important for selecting the right conditional branch. The remaining fields can always be accessed using the field selectors as and where needed.

5 Scope Rules

This is an exercise on understanding Haskell's scope rules. In the following code fragments, draw an arrow from each *use* of a variables (`x`, `y`, etc.) to its *defining occurrence*, if it has one in the provided fragment. For example for a code fragment like


```
let x = 3 in x + x
```

you would draw an arrow from each of the `x`'s in the expression `x + x` to the `x` in `x = 3` as that is the corresponding defining occurrence. Note that a particular variable *name*, like `x`, may be used for more than one variable even within the scope of a definition for the variable in question since inner definitions are allowed to *shadow* outer definitions in Haskell. For example, the value of the Haskell expression

```
let x = 7 in
  (let x = 3 in x + x) * x
```

is 42. Note that the following fragments are not examples of good style Haskell code! They have deliberately been made somewhat confusing to make a good exercise on Haskell's scope rules.

1.

```
f xs ys =
    let xs = x : xs in take 10 (ys ++ xs)
  where
    x = head xs
```
2.

```
f x y =
    let n = 3 in take n (g y) ++ take n (g x)
  where
    g x = take n xys
      where
        xys = x : yxs
        yxs = y : xys
    n = 10
```
3.

```
f xxs@(x:xs) =
    case xs of
      []      -> [x] : take n (repeat xs)
    (x:xs) -> [x] : take n (repeat xs)
  where
    n = length xxs
```